

Vector Quantization

This page contains information related to *vector quantization* (VQ). Currently this page includes information about VQ with regards to compression. In the future, we will make this page a more comprehensive VQ page.

In what applications is VQ used?

Vector quantization is used in many applications such as image and voice compression, voice recognition (in general statistical pattern recognition), and surprisingly enough in volume rendering (I have no idea how VQ is used in volume rendering!).

What is VQ?

A vector quantizer maps k -dimensional vectors in the vector space R^k into a finite set of vectors $Y = \{y_i: i = 1, 2, \dots, N\}$. Each vector y_i is called a code vector or a *codeword*. and the set of all the codewords is called a *codebook*. Associated with each codeword, y_i , is a nearest neighbor region called *Voronoi* region, and it is defined by:

$$V_i = \{x \in R^k : \|x - y_i\| \leq \|x - y_j\|, \text{ for all } j \neq i\}$$

The set of Voronoi regions partition the entire space R^k such that:

$$\bigcup_{i=1}^N V_i = R^k$$

$$\bigcap_{i=1}^N V_i = \phi \quad \text{for all } i \neq j$$

As an example we take vectors in the two dimensional case without loss of generality. Figure 1 shows some vectors in space. Associated with each cluster of vectors is a representative codeword. Each codeword resides in its own Voronoi region. These regions are separated with imaginary lines in figure 1 for illustration. Given an input vector, the codeword that is chosen to represent it is the one in the same Voronoi region.




Figure 1: Codewords in 2-dimensional space. Input vectors are marked with an x, codewords are marked with red circles, and the Voronoi regions are separated with boundary lines.

The representative codeword is determined to be the closest in Euclidean distance from the input vector. The Euclidean distance is defined by:

$$d(x, y_i) = \sqrt{\sum_{j=1}^k (x_j - y_{ij})^2}$$

where x_j is the j th component of the input vector, and y_{ij} is the j th component of the codeword y_i .

How does VQ work in compression?

A vector quantizer is composed of two operations. The first is the encoder, and the second is the decoder. The encoder takes an input vector and outputs the index of the codeword that offers the lowest distortion. In this case the lowest distortion is found by evaluating the Euclidean distance between the input vector and each codeword in the codebook. Once the closest codeword is found, the index of that codeword is sent through a channel (the channel could be a computer storage, communications channel, and so on). When the encoder receives the index of the codeword, it replaces the index with the associated codeword. Figure 2 shows a block diagram of the operation of the encoder and decoder.

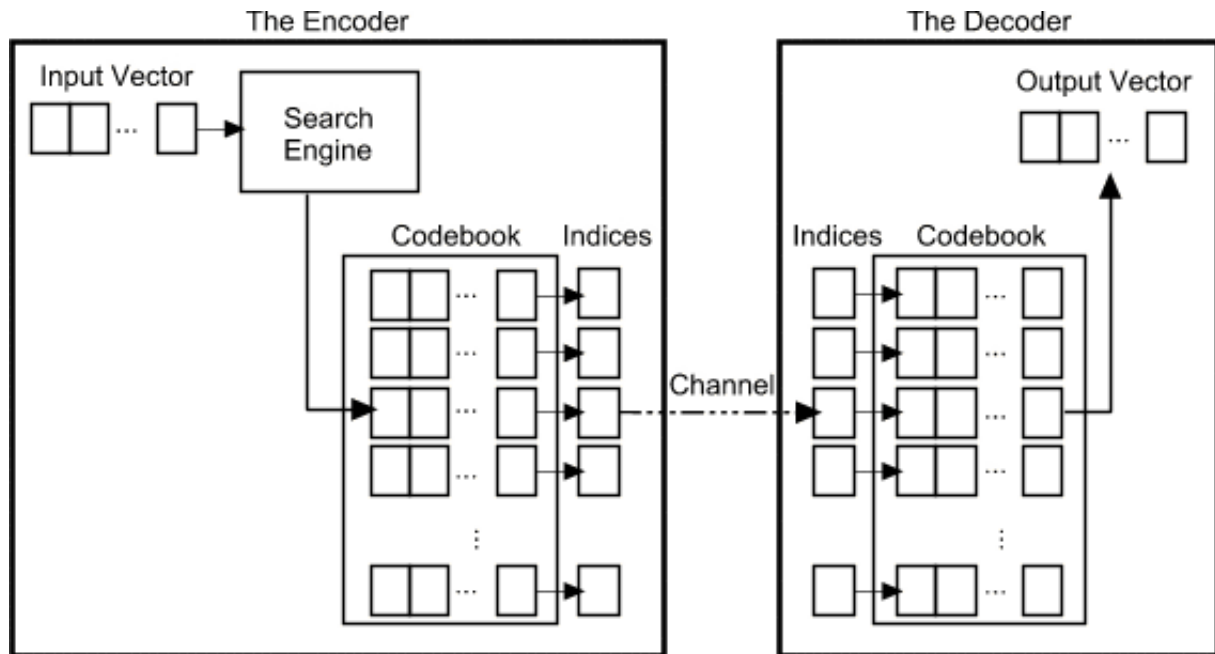


Figure 2: The Encoder and decoder in a vector quantizer. Given an input vector, the closest codeword is found and the index of the codeword is sent through the channel. The decoder receives the index of the codeword, and outputs the codeword.

How is the codebook designed?

So far we have talked about the way VQ works, but we haven't talked about how to generate the codebook. What code words best represent a given set of input vectors? How many should be chosen?

Unfortunately, designing a codebook that best represents the set of input vectors is NP-hard. That means that it requires an exhaustive search for the best possible codewords in space, and the search increases exponentially as the number of codewords increases (if you can find an optimal solution in polynomial time your name will go down in history forever). We therefore resort to suboptimal codebook design schemes, and the first one that comes to mind is the simplest. It is named [LBG](#) for Linde-Buzo-Gray, the authors of this idea. This algorithm is similar to the k-means algorithm.

The algorithm

1. **Determine the number of codewords, N , or the size of the codebook.**
2. **Select N codewords at random, and let that be the initial codebook.** The initial codewords can be randomly chosen from the set of input vectors.
3. **Using the Euclidean distance measure clusterize the vectors around each codeword.** This is done by taking each input vector and finding the Euclidean distance between it and each codeword. The input vector belongs to the cluster of the codeword that yields the minimum distance.
4. **Compute the new set of codewords.** This is done by obtaining the average of each cluster. Add the component of each vector and divide by the number of vectors in the cluster.

$$y_i = \frac{1}{m} \sum_{j=1}^m x_{ij}$$

where i is the component of each vector (x, y, z, \dots directions), m is the number of vectors in the

cluster.

1. **Repeat steps 2 and 3 until the either the codewords don't change or the change in the codewords is small.**

This algorithm is by far the most popular, and that is due to its simplicity. Although it is locally optimal, yet it is very slow. The reason it is slow is because for each iteration, determining each cluster requires that each input vector be compared with all the codewords in the codebook (We have programmed this algorithm in C, and for an 512x512 image, a codebook of 256, and vectors in 4 dimensions, the generation of the codebook took about 20 minutes on an HP machine).

There are many other methods to designing the codebook, methods such as [Pairwise Nearest Neighbor \(PNN\)](#), [Simulated Annealing](#), [Maximum Descent \(MD\)](#), and [Frequency-Sensitive Competitive Learning \(FSCL\)](#), etc.

How does the search engine work?

Although VQ offers more compression for the same distortion rate as scalar quantization and PCM, yet is not as widely implemented. This due to two things. The first is the time it takes to generate the codebook, and second is the speed of the search. [Many algorithms](#) have been proposed to increase the speed of the search. Some of them reduce the math used to determine the codeword that offers the minimum distortion, other algorithms preprocess the codewords and exploit underlying structure.

The simplest search method, which is also the slowest, is full search. In full search an input vector is compared with every codeword in the codebook. If there were M input vectors, N codewords, and each vector is in k dimensions, then the number of multiplies becomes kMN , the number of additions and subtractions become $MN((k - 1) + k) = MN(2k-1)$, and the number of comparisons becomes $MN(k - 1)$. This makes full search an expensive method.

What is the measure of performance VQ?

How does one rate the performance of a compressed image or sound using VQ? There is no good way to measure the performance of VQ. This is because the distortion that VQ incurs will be evaluated by us humans and that is a subjective measure. Don't despair! We can always resort to good old *Mean Squared Error* (MSE) and *Peak Signal to Noise Ratio* (PSNR). MSE is defined as follows:

$$MSE = \frac{1}{M} \sum_{i=1}^M (\hat{x}_i - x_i)^2$$

where M is the number of elements in the signal, or image. For example, if we wanted to find the MSE between the reconstructed and the original image, then we would take the difference between the two images pixel by pixel, square the results, and average the results.

The PSNR is defined as follows:

$$PSNR = 10 \log_{10} \left(\frac{(2^n - 1)^2}{MSE} \right)$$

where n is the number of bits per symbol. As an example, if we want to find the PSNR between two 256 gray level images, then we set n to 8 bits.

Some sites with VQ

Although there are many web pages on VQ, but the majority of them are not generalized. The majority of these sites describe new ways of implementing VQ or some of its applications. Unfortunately, most of these sites don't go into the detail of their work. I have therefore limited the links to those sites that contain useful information or tools.

- For a really brief description of VQ you can visit this [FAQ](#) sheet.
 - [Jim Fowler](#) was a graduate at The Ohio State University. Now he is teaching at the Department of Electrical & Computer Engineering at Mississippi State University. He has developed a wonderful package, [QccPack](#), for quantization, data compression and coding that includes VQ tools. He also worked with [video coding using VQ](#). He is a definite VQer.
 - [Possibilistic Clustering in Kohonen Networks for Vector Quantization](#). VQ using neural nets.
 - [Light Field Compression using Wavelet Transform and Vector Quantization](#).
 - [Robert Gray](#) teaches at Stanford University, and within his class of [Quantization and Data compression](#) he devotes a topic to [vector quantization](#). This is an acrobat file of slides.
 - [Vivek Goyal](#) has some interesting papers on VQ. Check out his [publications](#) page.
 - Another Neural Network based VQ with code called, [Predictive Residual Vector Quantization \(PRVQ\) CODEC](#).
 - [Dynamic Learning Vector Quantization \(DLVQ\)](#).
-

VQers

This is a list of people, in alphabetical order by last name, who constantly work on VQ. I call them VQers. They are a must know! If you think you are VQer, then send me an [email](#), and I will include your name in the list.

- [Stanley Ahalt](#).
 - [Jim Fowler](#).
 - [Allen Gersho](#).
 - [Robert M. Gray](#).
 - [Batuhan Ulug](#).
-

You can email me at mohamedqasem@yahoo.com if you have any questions or comments.
[back to my home page](#)